

Viewstamped Replication, Latency vs Throughput, Multiple Shards

